

# COMPSCI 574: Intelligent Visual Computing Spring 23

Prachi Jain

## Assignment 2: Multi-View CNN for Shape Classification

### 1 Task A [70%]

Change starter code in "model.py" to define specified convnet:

#### 1.1 Implementation

Please find attached code in *model.py* for the design of convnet.

```
1 import torch
2 import torch.nn as nn
3
4 #Utility function for calculating output size of previous convolution layer
  which can be used as a parameter
5 #for normalized_shape argument in nn.LayerNorm - "think about what the parameter
  normalized_shape should be"
6 import numpy as np
7 def conv2d_output_size(input_size, out_channels, padding, kernel_size, stride,
  dilation=None):
8     """According to https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html
9     """
10     if dilation is None:
11         dilation = (1, ) * 2
12     if isinstance(padding, int):
13         padding = (padding, ) * 2
14     if isinstance(kernel_size, int):
15         kernel_size = (kernel_size, ) * 2
16     if isinstance(stride, int):
17         stride = (stride, ) * 2
18
19     output_size = (
20         out_channels,
21         np.floor((input_size[0] + 2 * padding[0] - dilation[0] * (kernel_size[0]
22         - 1) - 1) / stride[0] + 1).astype(int),
```

```

23         np.floor((input_size[1] + 2 * padding[1] - dilation[1] * (kernel_size[1]
24         - 1) - 1) /
25         stride[1] + 1).astype(int)
26     )
27     return output_size
28 class CNN(nn.Module):
29
30     def __init__(self, num_classes):
31         super(CNN, self).__init__()
32         """
33
34         DEFINE YOUR NETWORK HERE
35
36         """
37
38         self.num_classes = num_classes
39
40         # Convolution layers self.conv1, self.depthconv1, self.depthconv2
41         # followed by norm layers and leaky relu
42         # would not have biases since normalisation layer will re-center the
43         # data anyway, removing the bias and
44         # making it a useless trainable parameter.
45
46         #1 Convolution layer with 8 filters applying 7x7 filters on input image.
47         # Set stride to 1 and padding to '0'.
48         self.conv1 = nn.Conv2d(in_channels = 1, out_channels = 8, kernel_size =
49         7, stride = 1, padding = 0, bias = False)
50         #H,W of Input: 112, 112
51         outputSize1 = conv2d_output_size([112, 112], out_channels = 8,
52         kernel_size = 7, stride = 1, padding = 0)
53
54         #2. Normalization layer
55         self.norm1 = nn.LayerNorm(normalized_shape = [*outputSize1]) # <---
56         # Normalize activations over C, H, and W shape of the Convolution Output
57
58         #3, 7, 12 Leaky ReLU layer with 0.01 'leak' for negative input
59         self.leakyReLU = nn.LeakyReLU(0.1)
60
61         #4, 8, 13 Max pooling layer operating on 2x2 windows with stride 2
62         self.MaxPool = nn.MaxPool2d(kernel_size = 2, stride = 2)
63         outputSize2 = conv2d_output_size(outputSize1[1:], out_channels = 8,
64         kernel_size = 2, stride = 2, padding = 0)
65
66         #5. Depthwise convolution layer with 8 filters applying 7x7 filters on
67         # the feature map of the previous layer.
68         # Set stride to 2 and padding to 0. Groups should be 8 to enable
69         # depthwise convolution.
70         self.depthconv1 = nn.Conv2d(in_channels = 8, out_channels = 8,
71         kernel_size = 7, stride = 2, padding = 0, groups = 8, bias = False)
72         outputSize3 = conv2d_output_size(outputSize2[1:], out_channels = 8,
73         kernel_size = 7, stride = 2, padding = 0)
74
75         #Max Pool outputsize calculation
76         outputSize4 = conv2d_output_size(outputSize3[1:], out_channels = 8,
77         kernel_size = 2, stride = 2, padding = 0)
78
79         #6. Normalization layer

```

```

68         self.norm2 = nn.LayerNorm(normalized_shape = [*outputSize3])
69
70         #9. Pointwise convolution layer with 16 filters applying 1x1 filters on
the feature map of the previous layer
71         #(stride is 1, and no padding).
72         self.pointconv1 = nn.Conv2d(in_channels = 8, out_channels = 16,
kernel_size = 1, stride = 1, padding = 0, bias = True)
73         outputSize5 = conv2d_output_size(outputSize4[1:], out_channels = 16,
kernel_size = 1, stride = 1, padding = 0)
74
75         #10. Depthwise convolution layer with 16 filters applying 7x7 filters on
the feature map of the previous layer.
76         #Set stride to 1 and padding to 0. Groups should be 16 to enable
depthwise convolution
77         self.depthconv2 = nn.Conv2d(in_channels = 16, out_channels = 16,
kernel_size = 7, stride = 1, padding = 0, groups = 16, bias = False)
78         outputSize6 = conv2d_output_size(outputSize5[1:], out_channels = 16,
kernel_size = 7, stride = 1, padding = 0)
79
80         #11. Normalization layer
81         self.norm3 = nn.LayerNorm(normalized_shape = [*outputSize6])
82
83         #14. Pointwise convolution layer with 32 filters applying 1x1 filters on
the feature map of the previous layer
84         #(stride is 1, and no padding).
85         self.pointconv2 = nn.Conv2d(in_channels = 16, out_channels = 32,
kernel_size = 1, stride = 1, padding = 0, bias = True)
86
87         #15. Fully connected layer with K outputs, where K is the number of
categories. Include the bias.
88         #Implement the fully connected layer as a convolutional one (think how
and why).
89         self.fullconnect = nn.Conv2d(in_channels = 32, out_channels =
num_classes, kernel_size = 3, bias = True)
90
91         self.apply(self._init_weights)
92
93     def _init_weights(self, module):
94         """ Initialize the weights """
95         if isinstance(module, nn.Conv2d):
96             # Initialize weights according to the Kaiming Xe's uniform
distribution scheme
97             # for conv layers followed by leaky ReLU activations
98             nn.init.kaiming_uniform_(self.conv1.weight)
99             nn.init.kaiming_uniform_(self.depthconv1.weight)
100             nn.init.kaiming_uniform_(self.depthconv2.weight)
101
102             # Initialize weights according to the Xavier uniform distribution
scheme for pointwise convolution layers
103             # and fully connected layer (i.e., layers not followed by ReLUs).
104             nn.init.xavier_uniform_(self.pointconv1.weight)
105             nn.init.xavier_uniform_(self.pointconv2.weight)
106             nn.init.xavier_uniform_(self.fullconnect.weight)
107
108             # Initialize any biases to 0.
109             if module.bias is not None:
110                 module.bias.data.zero_()
111

```

```

112 def forward(self, x):
113     """
114
115     DEFINE YOUR FORWARD PASS HERE
116
117     """
118     h1 = self.conv1(x)
119     h2 = self.norm1(h1)
120     h3 = self.leakyReLU(h2)
121     h4 = self.MaxPool(h3)
122
123     #Combination of depthwise and pointwise convolution produces a block
    called "Depthwise Separable Convolution"
124     h5 = self.depthconv1(h4)
125     h6 = self.norm2(h5)
126     h7 = self.leakyReLU(h6)
127     h8 = self.MaxPool(h7)
128     h9 = self.pointconv1(h8)
129
130     #Combination of depthwise and pointwise convolution produces a block
    called "Depthwise Separable Convolution"
131     h10 = self.depthconv2(h9)
132     h11 = self.norm3(h10)
133     h12 = self.leakyReLU(h11)
134     h13 = self.MaxPool(h12)
135     h14 = self.pointconv2(h13)
136
137     out = self.fullconnect(h14)
138     return out

```

## 2 Task B [30%]

Modify the function of "testMVImageClassifier.py" such that the function outputs category predictions per shape and test error averaged over all the test shapes with two strategies:

- Mean view-pooling
- Mean view-pooling

### 2.1 Implementation

Please find attached code of *testMVImageClassifier.py*.

```

1 import os
2 import numpy as np
3 import torch
4 from torch.autograd import Variable
5 from data_utils import grayscale_img_load, listdir
6

```

```

7 def testMVImageClassifier(dataset_path, model, info, pooling = 'mean', cuda=
  False, verbose=False):
8
9     # save pytorch model to eval mode
10    model.eval()
11    if (cuda):
12        model.cuda()
13
14    test_err = 0
15    count = 0
16    print("=>Testing...")
17
18    # for each category
19    for idx, c in enumerate(info['category_names']):
20        category_full_dir = os.path.join(dataset_path, c)
21        shape_dirs        = listdir(category_full_dir)
22        print('=>Loading shape data: %s'%(c))
23
24        # for each shape
25        for s in shape_dirs:
26            if verbose: print('=>Loading shape data: %s %s'%(s, c))
27            views = listdir(os.path.join(category_full_dir, s))
28            scores = np.zeros((len(views), len(info['category_names'])))
29            count += 1
30
31            # for each view
32            for i, v in enumerate(views):
33                image_full_filename = os.path.join(category_full_dir, s, v)
34                if 'png' not in image_full_filename : continue
35                if verbose: print(' => Loading image: %s ...'%
image_full_filename)
36                im = grayscale_img_load(image_full_filename)/255.
37                im -= info['data_mean']
38                im = Variable(torch.from_numpy(im.astype('float32')),
requires_grad=False).unsqueeze(0)
39                # get predicted scores for each view
40                if (cuda):
41                    im = im.cuda()
42                    scores[i, :] = model(im).detach().cpu().numpy().squeeze()
43                else:
44                    scores[i, :] = model(im).detach().numpy().squeeze()
45
46                '''
47                YOUR CODE GOES HERE
48                1) Get category predictions per shape and test error averaged over
all the test shapes.
49                2) Implement 2 strategies: 1) mean and 2) max view-pooling by
specifying input arg 'pooling', like
50                >> pooling = 'mean' or pooling = 'max'
51
52                '''
53                #predicted_label = 0 # obviously change this
54
55                # Mean view-pooling
56                if pooling == 'mean':
57                    predicted_label = np.argmax(np.mean(scores, axis = 0))
58                # Max view-pooling
59                elif pooling == 'max':

```

```

60     predicted_label = np.argmax(np.max(scores, axis = 0))
61
62     if predicted_label != idx:
63         test_err += 1
64
65     if verbose: print('predicted label: %s, ground-truth label: %s\n'%(
66 info['category_names'][predicted_label] ,c))
67
68     test_err = test_err / count
69     print('Test error: %f%%\n'%(test_err * 100))

```

### 3 Performance Evaluation

- Test Error for Mean View Pooling = **14.000000%**
- Test Error for Max View Pooling = **17.000000%**




Figure 1: Plot showing a). Epoch vs Training and Validation loss, b) 2) Epoch vs Training and Validation Top1 Error and c). 2) Epoch vs Training and Validation Top1 Error